ZHUOFENG LI

zhuofeng-li.github.io | zhuofengli12345@gmail.com
Zhuofeng-Li (5K+ Stars) | Zhuofeng Li

EDUCATION

Stanford University, 450 Jane Stanford Way Stanford, CA 94305–2004 *June, 2025 – Dec, 2025*

Visiting Student

Advisor: Prof. Yejin Choi and Prof. James Zou

Texas A&M University, 1226 TAMU, College Station, TX 77843 *2025 – 2027*

Ph.D. in Computer Science

Advisor: Prof. Yu Zhang

INDUSTRY EXPERIENCE

Google DeepMind *Sep, 2026*

Student Researcher (Self-Evolving, Post-training)

Sea AI Lab & University of Waterloo *Feb, 2025 – Aug, 2025*

Research Intern (Agentic Reinforcement Learning)

RESEARCH INTERESTS

Fields Self-Evolving, Agentic System, Agentic RL, LLM Post-training,

Proficient Slime, Verl, VLLM, Sglang, Ray, Megatron, DeepSpeed, Harbor

RESEARCH

* indicates equal contribution

Highlight

1. **OpenRSI Index.** [website] [code]

   RSI

   Yuetai Li*, **Zhuofeng Li* (Tech Lead)**, Yichen Feng*, Xinyang Han*, Wenhao Chai*, Chenglei Si*, Shangding Gu, Rilyn Han, Zhengyu Hu, Hang Hua, Yue Huang, Cara Li, Hanchen Li, Shujia Liang, Xiang Liu, Pan Lu, Bohan Lyu, Zixian Ma, Qiuyang Mang, Rulin Shao, Venkat Srinivasan, Yiyou Sun, Guan Wang, Ning Wang, Jian Xie, Brian Yu, Gaoyuan Zhang, Qizheng Zhang, Weichen Zhang, Wenhu Chen, **Yejin Choi**, **Jianfeng Gao**, Hannaneh Hajishirzi, Pang Wei Koh, **Ranjay Krishna**, Karthik Narasimhan, Kunle Olukotun, Rameswar Panda, Radha Poovendran, **Dawn Song**, Yu Su, **Diyi Yang**, Ming-Hsuan Yang, Xiang Yue, **Jian Zhang**, Yu Zhang, Luke Zettlemoyer

   Scaling RSI on Thousands of GPUs

2. **In-The-Flow Agentic System Optimization.** [arxiv] [code]

   Agentic RL | Multi-Agent | Training Infra

   Zhuofeng Li*, Haoxiang Zhang*, Seungju Han, Sheng Liu, Jianwen Xie, **Yu Zhang**, **Yejin Choi**, **James Zou**, **Pan Lu**

   ICLR 2026 Oral (Top 1.1%) & 2K+ Stars

3. **VerlTool: A unified and easy-to-extend tool-agent training framework.** [arxiv] [code]

   Agentic RL | Agentic Coding | Agentic Search | Training Infra

   Donfu Jiang*, **Zhuofeng Li***, Yi Lu*, Zhiheng Lv, Ping Nie, Chao Du, Tianyu Pang, **Wenhu Chen**

   ICLR 2026 SPOT Best Paper (Top 1) & 1K+ Stars

4. **Beyond Semantic Similarity: Rethinking Retrieval for Agentic Search via Direct Corpus Interaction.** [arXiv] [code]

   Agentic Search | Position Paper

   Zhuofeng Li, Haoxiang Zhang, Cong Wei, Pan Lu, Ping Nie, Yi Lu, Yuyang Bai, Shangbin Feng, Hangxiao Zhu, Ming Zhong, Yuyu Zhang, Jianwen Xie, **Yejin Choi**, **James Zou**, **Jiawei Han**, **Wenhu Chen**, **Jimmy Lin**, **Dongfu Jiang**, **Yu Zhang**

   NeurIPS 2026 Spotlight

5. **OpenResearcher: A Fully Open Pipeline for Long-Horizon Deep Research Trajectory Synthesis.** [arxiv] [code]

   Agentic Search | Synthetic Data Generation

   Zhuofeng Li*, Dongfu Jiang*, Xueguang Ma, Haoxiang Zhang, Ping Nie, Yuyu Zhang, Kai Zou, Jianwen Xie, **Yu Zhang**, **Wenhu Chen**

   EMNLP 2026 & Adopted by NVIDIA's Nemotron models and Amazon Rufus & Top 3 Trending Datasets on Hugging Face & 1.2K+ Stars

Selected Publications

1. **Dr-DCI: Scaling Direct Corpus Interaction via Dynamic Workspace Expansion.**

   Yi Lu*, **Zhuofeng Li***, Ping Nie, Haoxiang Zhang, Yuyu Zhang, Kai Zou, Wenhu Chen, Jimmy Lin, Dongfu Jiang, Yu Zhang

   In a submission to ICLR 2026 [arXiv]

2. **Self-Retrospection Distillation: Turning Post-hoc Experiences into Prior Foresight.**

   Haoxiang Zhang*, Qinglin Chen*, Hiroaki Hayashi*, **Zhuofeng Li***, Siming Zhang, Jiaxin Zhang, Jixuan Chen, Fang Wu, Pan Lu, Silvio Savarese, Julian McAuley, Chien-Sheng Wu

   In a submission to ICLR 2026 [arXiv]

3. **Masking Stale Observations Helps Search Agents – Until It Doesn't: A Regime Map and Its Mechanisms.**

   Haoxiang Zhang*, Qixin Xu*, **Zhuofeng Li***, Lei Zhang, Pengcheng Jiang, Yu Zhang, Julian McAuley

   EMNLP 2026 [arXiv]

4. **ClawBench: Can AI Agents Complete Everyday Online Tasks?**

   Core Contributor

   EMNLP 2026 & **Adopted by Meta Superintelligence Labs** [arXiv]

5. **ReviewGrounder: Improving Review Substantiveness with Rubric-Guided, Tool-Integrated Agents.**

   Zhuofeng Li, Yi Lu, Dongfu Jiang, Haoxiang Zhang, Yuyang Bai, Chuan Li, Yu Wang, Shuiwang Ji, Jianwen Xie, Yu Zhang

ACL 2026 [arXiv]

6. **GReF: A Unified Generative Framework for Efficient Reranking via Ordered Multi-token Prediction.**

   Zhijie Lin*, **Zhuofeng Li***, ChengLei Dai, Wentian Bao, Shuai Lin, Yun En Yu, Haoxiang Zhang, Liang Zhao

   CIKM 2025 [arXiv]

7. **TEG-DB: A Comprehensive Dataset and Benchmark of Textual-Edge Graphs.**

   Zhuofeng Li, Zixing Gou, Xiangnan Zhang, Zhongyuan Liu, Sirui Li, Yuntong Hu, Chen Ling, Zheng Zhang, Liang Zhao.

   NeurIPS 2024 [arXiv]

RESEARCH EXPERIENCE

Amazon, Palo Alto, CA
Foundation Model Team Applied Scientist Intern *May, 2026 – Aug, 2026*
Project: Large-Scale Post-Training and Agentic RL

- Conducted **large-scale post-training and reinforcement learning** for foundation models, developing and scaling distributed training pipelines for frontier model development.
- Trained and optimized **Search Agents, Coding Agents, and Terminal Agents** with large-scale agentic RL, enabling models to perform long-horizon reasoning, tool use, and environment interaction.

Stanford University, 450 Jane Stanford Way Stanford, CA 94305–2004
Department of Computer Science Zou's Group Choi's xlab *June, 2025 – Dec, 2025*
Research Assistant, Advisor: Prof. James Zou and Prof. Yejin Choi
Project: In-the-Flow Agentic System Optimization [arxiv] [website]

- Introduce AgentFlow, a trainable tool-integrated agentic system to overcome the scalability and generalization limits of today's tool-augmented reasoning approaches.
- Pioneer a novel RL paradigm to directly optimize the agent within the system in an online fashion.
- Achieve strong performance on 10 benchmarks with a 7B backbone: +14.9% search, +14.0% agentic, +14.5% math, +4.1% scientific tasks; **surpassing larger-scale models like GPT**.

University of Waterloo, 200 University Ave. West, Waterloo, Ontario
Department of Computer Science TIGER-AI-Lab *January, 2025 – June, 2025*
Research Assistant, Advisor: Prof. Wenhu Chen
Project: Agentic Tool-Use LLMs through RL [arXiv] [code]

- Propose a novel agentic async tool-use RL training framework.
- Achieve strong performance across diverse benchmarks, including search, math and swe tasks.