

ZHUOFENG LI



[zhuofeng-li.github.io](https://github.com/zhuofeng-li)
[Zhuofeng-Li](#) (3K+ Stars)



zhuofengli12345@gmail.com
[Zhuofeng Li](#)

EDUCATION

Stanford University, 450 Jane Stanford Way Stanford, CA 94305–2004 *June, 2025 – Dec, 2025*

Visiting Student

Advisor: Prof. [Yejin Choi](#) and Prof. [James Zou](#)

Texas A&M University, 1226 TAMU, College Station, TX 77843 *Aug, 2025 – Aug, 2028 (expected)*

Ph.D. in Computer Science

Advisor: Prof. [Yu Zhang](#)

INDUSTRY EXPERIENCE

Amazon Rufus *May, 2026 – Aug, 2026*

Foundation Model Team Research Intern (Agent Swarm RL) Mentor: [Dr. Qingyu Yin](#) and [Dr. Xin Liu](#)

Sea AI Lab & University of Waterloo *Feb, 2025 – Aug, 2025*

Research Intern (Agentic Reinforcement Learning) Mentor: [Prof. Wenhui Chen](#)

RESEARCH INTERESTS

Fields Agentic System, Reasoning LLMs/VLMs, Reinforcement Learning, Multi-modal LLM

Proficient Slime, Verl, VLLM, Sglang, Ray, Megatron, DeepSpeed, Harbor

RESEARCH

* indicates equal contribution

Highlight

1. **In-The-Flow Agentic System Optimization.** [\[arxiv\]](#) [\[code\]](#)

[Agentic RL](#) [Multi-Agent](#) [Training Infra](#)

Zhuofeng Li*, Haoxiang Zhang*, Seungju Han, Sheng Liu, Jianwen Xie, **Yu Zhang**, **Yejin Choi**, James Zou, Pan Lu

ICLR 2026 Oral (Top 1.1%) & 1.8K+ Stars

2. **VerlTool: A unified and easy-to-extend tool-agent training framework.** [\[arxiv\]](#) [\[code\]](#)

[Agentic RL](#) [Agentic Coding](#) [Agentic Search](#) [Training Infra](#)

Donfu Jiang*, **Zhuofeng Li***, Yi Lu*, Zhiheng Lv, Ping Nie, Chao Du, Tianyu Pang, **Wenhui Chen**

ICLR 2026 SPOT Best Paper (Top 1) & 900+ Stars

3. **Beyond Semantic Similarity: Rethinking Retrieval for Agentic Search via Direct Corpus Interaction.** [\[arXiv\]](#) [\[code\]](#)

[Agentic Search](#) [Position Paper](#)

Zhuofeng Li, Haoxiang Zhang, Cong Wei, Pan Lu, Ping Nie, Yi Lu, Yuyang Bai, Shangbin Feng, Hangxiao Zhu, Ming Zhong, Yuyu Zhang, Jianwen Xie, **Yejin Choi**, **James Zou**, **Jiawei Han**, **Wenhu Chen**, **Jimmy Lin**, **Dongfu Jiang**, **Yu Zhang**

#1 paper of the day on Hugging Face

4. **OpenResearcher: A Fully Open Pipeline for Long-Horizon Deep Research Trajectory Synthesis.** [\[arxiv\]](#) [\[code\]](#)

Agentic Search

Synthetic Data Generation

Zhuofeng Li*, Dongfu Jiang*, Xueguang Ma, Haoxiang Zhang, Ping Nie, Yuyu Zhang, Kai Zou, Jianwen Xie, **Yu Zhang**, **Wenhu Chen**

Adopted by NVIDIA's Nemotron models & Top 3 Trending Datasets on Hugging Face & 700+ Stars

Selected Publications

1. **In-The-Flow Agentic System Optimization for Effective Planning and Tool Use.** [\[arxiv\]](#) [\[code\]](#)

Zhuofeng Li*, Haoxiang Zhang*, Seungju Han, Sheng Liu, Jianwen Xie, **Yu Zhang**, **Yejin Choi**, **James Zou**, **Pan Lu**

ICLR 2026 Oral (Top 1.1%) & NeurIPS 2025 Workshop Best Paper Runner-up

2. **ReviewGrounder: Improving Review Substantiveness with Rubric-Guided, Tool-Integrated Agents.** *ACL 2026.* [\[arXiv\]](#)

Zhuofeng Li, Yi Lu, Dongfu Jiang, Haoxiang Zhang, Yuyang Bai, Chuan Li, Yu Wang, Shuiwang Ji, Jianwen Xie, **Yu Zhang**

3. **Masking Stale Observations Helps Search Agents – Until It Doesn't: A Regime Map and Its Mechanisms.** [\[arXiv\]](#)

Haoxiang Zhang*, Qixin Xu*, **Zhuofeng Li***, Lei Zhang, Pengcheng Jiang, Yu Zhang, Julian McAuley

4. **ClawBench: Can AI Agents Complete Everyday Online Tasks?**

Adopted by Meta Superintelligence Labs

5. **ImagenWorld: Stress-Testing Image Generation Models with Explainable Human Evaluation on Open-ended Real-World Tasks.** *ICLR 2026.*

6. **VideoEval-Pro: Robust and Realistic Long Video Understanding Evaluation.** *ICLR 2026.*

7. **VideoScore2: Think before You Reward in Video Generation.** *TMLR 2026.*

8. **StructEval: Benchmarking LLMs' Capabilities to Generate Structural Outputs.** *TMLR 2025. (J2C Certification at ICML/NeurIPS/ICLR)* [\[arXiv\]](#)

9. **GReF: A Unified Generative Framework for Efficient Reranking via Ordered Multi-token Prediction.** *CIKM 2025.* [\[arXiv\]](#)

Zhijie Lin*, **Zhuofeng Li***, ChengLei Dai, Wentian Bao, Shuai Lin, Yun En Yu, Haoxiang Zhang, **Liang Zhao**

10. **Avoiding Structural Pitfalls: Self-Supervised Low-Rank Feature Tuning for Graph Test-Time Adaptation.** *TMLR 2025.* [\[arXiv\]](#)

Haoxiang Zhang*, **Zhuofeng Li***, Qiannan Zhang, Ziyi Kou, Juncheng Li, Shichao Pei.

11. **TEG-DB: A Comprehensive Dataset and Benchmark of Textual-Edge Graphs.** *NeurIPS 2024*. [\[arXiv\]](#)[\[code\]](#)
Zhuofeng Li*, Zixing Gou*, Xiangnan Zhang, Zhongyuan Liu, Sirui Li, Yuntong Hu, Chen Ling, Zheng Zhang, **Liang Zhao**.
12. **Learning from Novel Knowledge: Continual Few-shot Knowledge Graph Completion.** *CIKM 2024*. [\[arXiv\]](#)
Zhuofeng Li*, Haoxiang Zhang*, Qiannan Zhang, Ziyi Kou, Shichao Pei.

RESEARCH EXPERIENCE

Amazon Rufus, 611 Cowper St, Palo Alto, CA 94301, USA

Foundation Model Team

May, 2026 – Aug, 2026

Applied Scientist Intern, Advisor: Dr. [Qingyu Yin](#) and Dr. [Xin Liu](#)

Project: Large-scale post-training and agent swarm RL

Stanford University, 450 Jane Stanford Way Stanford, CA 94305-2004

Department of Computer Science [Zou's Group](#) [Choi's xlab](#)

June, 2025 – Dec, 2025

Research Assistant, Advisor: Prof. [James Zou](#) and Prof. [Yejin Choi](#)

Project: In-the-Flow Agentic System Optimization [\[arxiv\]](#) [\[website\]](#)

- Introduce AgentFlow, a trainable tool-integrated agentic system to overcome the scalability and generalization limits of today's tool-augmented reasoning approaches.
- Pioneer a novel RL paradigm to directly optimize the agent within the system in an online fashion.
- Achieve strong performance on 10 benchmarks with a 7B backbone: +14.9% search, +14.0% agentic, +14.5% math, +4.1% scientific tasks; **surpassing larger-scale models like GPT**.

University of Waterloo, 200 University Ave. West, Waterloo, Ontario

Department of Computer Science [TIGER-AI-Lab](#)

January, 2025 – June, 2025

Research Assistant, Advisor: Prof. [Wenhu Chen](#)

Project: Agentic Tool-Use LLMs through RL [\[arXiv\]](#) [\[code\]](#)

- Propose a novel agentic async tool-use RL training framework.
- Achieve strong performance across diverse benchmarks, including math and search tasks.
- **Open-source tool-agent training framework [Verl-Tool](#) (900+ stars now).**